

# Bios 6301: HW 2 Jordan Clark

*Jordan Clark*

(informally) Due Thursday, 17 September, 1:00 PM

50 points total.

This assignment won't be submitted until we've covered Rmarkdown. Create R chunks for each question and insert your R code appropriately. Check your output by using the `Knit PDF` button in RStudio.

1. **Working with data** In the `datasets` folder on the course GitHub repo, you will find a file called `cancer.csv`, which is a dataset in comma-separated values (csv) format. This is a large cancer incidence dataset that summarizes the incidence of different cancers for various subgroups. (18 points)

1. Load the data set into R and make it a data frame called `cancer.df`. (2 points)

```
setwd("~/Desktop/Bios6301")
getwd()
```

```
## [1] "/Users/Kasasa/Desktop/Bios6301"
```

```
cancer.df<-read.csv("../datasets/cancer.csv")
head(cancer.df)
```

```
##   year      site      state  sex    race mortality
## 1 1999 Brain and Other Nervous System alabama Female   Black      0.00
## 2 1999 Brain and Other Nervous System alabama Female Hispanic 0.00
## 3 1999 Brain and Other Nervous System alabama Female   White  83.67
## 4 1999 Brain and Other Nervous System alabama  Male   Black      0.00
## 5 1999 Brain and Other Nervous System alabama  Male Hispanic 0.00
## 6 1999 Brain and Other Nervous System alabama  Male   White 103.66
## incidence population
## 1      19      623475
## 2       0      28101
## 3     110     1640665
## 4      18     539198
## 5       0      37082
## 6     145     1570643
```

```
names(cancer.df)
```

```
## [1] "year"      "site"      "state"     "sex"       "race"
## [6] "mortality" "incidence" "population"
```

2. Determine the number of rows and columns in the data frame. (2)

```
dim(cancer.df)
```

```
## [1] 42120      8
```

*#8 columns and 42120 rows*

3. Extract the names of the columns in `cancer.df`. (2)

```
col.names<-c(colnames(cancer.df))
col.names
```

```
## [1] "year"      "site"      "state"     "sex"       "race"
## [6] "mortality" "incidence" "population"
```

4. Report the value of the 3000th row in column 6. (2)

```
print(cancer.df[3000,6])
```

```
## [1] 350.69
```

5. Report the contents of the 172nd row. (2)

```
print(cancer.df[172,])
```

```
##      year                site state sex race mortality
## 172 1999 Brain and Other Nervous System nevada Male Black      0
##      incidence population
## 172          0          73172
```

6. Create a new column that is the incidence \*rate\* (per 100,000) for each row.(3)

```
cancer.df<-transform(cancer.df, i.rate=incidence/100000 )
head(cancer.df)
```

```
##   year                site state sex   race mortality
## 1 1999 Brain and Other Nervous System alabama Female   Black      0.00
## 2 1999 Brain and Other Nervous System alabama Female Hispanic 0.00
## 3 1999 Brain and Other Nervous System alabama Female   White  83.67
## 4 1999 Brain and Other Nervous System alabama Male     Black   0.00
## 5 1999 Brain and Other Nervous System alabama Male Hispanic 0.00
## 6 1999 Brain and Other Nervous System alabama Male     White 103.66
##   incidence population i.rate
## 1         19        623475 0.00019
## 2          0         28101 0.00000
## 3        110       1640665 0.00110
## 4         18        539198 0.00018
## 5          0         37082 0.00000
## 6        145       1570643 0.00145
```

7. How many subgroups (rows) have a zero incidence rate? (2)

```
zero<-0.0
zero
```

```
## [1] 0
```

```
cancer0.df<-cancer.df[cancer.df$i.rate %in% zero,]
head(cancer0.df)
```

```
##   year      site      state  sex   race mortality
## 2  1999 Brain and Other Nervous System alabama Female Hispanic      0
## 5  1999 Brain and Other Nervous System alabama  Male Hispanic      0
## 7  1999 Brain and Other Nervous System alaska  Female   Black      0
## 8  1999 Brain and Other Nervous System alaska  Female Hispanic      0
## 9  1999 Brain and Other Nervous System alaska  Female   White      0
## 10 1999 Brain and Other Nervous System alaska  Male    Black      0
##   incidence population i.rate
## 2         0        28101      0
## 5         0        37082      0
## 7         0        12710      0
## 8         0        11664      0
## 9         0       220036      0
## 10        0        13900      0
```

8. Find the subgroup with the highest incidence rate.(3)

```
which.max(cancer.df[, 'i.rate'])
```

```
## [1] 21387
```

```
head(order(cancer.df[, 'i.rate'], decreasing=T))
```

```
## [1] 21387 14367 327 7347 26847 28407
```

```
#Looks like row 21387
cancer.df[21387,]
```

```
##      year  site      state  sex  race mortality incidence population
## 21387 2002 Breast california Female White  3463.74    18774  13690681
##      i.rate
## 21387 0.18774
```

## 2. Data types (10 points)

1. Create the following vector: `x <- c("5","12","7")`. Which of the following commands will produce an error message? For each command, Either explain why they should be errors, or explain the non-erroneous result. (4 points)

```
x <- c("5","12","7")
class(x)
```

```
## [1] "character"
```

```
max(x)
```

```
## [1] "7"
```

```
#It returns a 7 because the elements are characters. 12 is not the max...  
#... because it starts with 1. So max returns a pseudo-alphabetic order.  
sort(x)
```

```
## [1] "12" "5" "7"
```

```
#Because our elements are characters, '1' is like 'a', '5' is like 'e' ...  
#... and '7' is like 'g', and they are ordered pseudo-alphabeticly.  
#(sum(x)) #I'm commenting this out because it won't compile otherwise.  
#It is an error because you cannot sum characters like you can numerically.
```

2. For the next two commands, either explain their results, or why they should produce errors. (3 points)

```
y <- c("5",7,12)  
class(y[2])
```

```
## [1] "character"
```

```
# y[2] + y[3]  
#It appears that the elements of this vector have taken the class...  
#...of the first element, which is a character.
```

3. For the next two commands, either explain their results, or why they should produce errors. (3 points)

```
z <- data.frame(z1="5",z2=7,z3=12)  
z[1,2] + z[1,3]
```

```
## [1] 19
```

```
#It appears that this worked because it's a data frame, rather than a vector.
```

3. **Data structures** Give R expressions that return the following matrices and vectors (*i.e.* do not construct them manually). (3 points each, 12 total)

1. (1, 2, 3, 4, 5, 6, 7, 8, 7, 6, 5, 4, 3, 2, 1)

```
xx<-(1:8)  
yy<-rev(xx)  
zz<-append(xx,yy)  
zz<-zz[-8]  
zz
```

```
## [1] 1 2 3 4 5 6 7 8 7 6 5 4 3 2 1
```

```
#somewhat manually constructed?
```

2. (1,2,2,3,3,3,4,4,4,4,5,5,5,5,5)

```
rep(1:5, 1:5)
```

```
## [1] 1 2 2 3 3 3 4 4 4 4 5 5 5 5 5
```

3.  $\begin{pmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix}$

```
0 & 1 & 1 \\
1 & 0 & 1 \\
1 & 1 & 0 \\
\end{pmatrix}
```

```
g<-c(0,1)
f<-c(0,1)
h<-c(1,0)
i<-expand.grid(g,f,h)
i
```

```
##   Var1 Var2 Var3
## 1    0    0    1
## 2    1    0    1
## 3    0    1    1
## 4    1    1    1
## 5    0    0    0
## 6    1    0    0
## 7    0    1    0
## 8    1    1    0
```

```
#This is obviously too much.
i.i<-i[which(rowSums(i)==2),]
i.i
```

```
##   Var1 Var2 Var3
## 2    1    0    1
## 3    0    1    1
## 8    1    1    0
```

```
#Just reorder this puppy
i.ii<-i.i[,c(2,1,3)]
i.ii
```

```
##   Var2 Var1 Var3
## 2    0    1    1
## 3    1    0    1
## 8    1    1    0
```

```
#Meandering solution but it worked.
```

```
4.  $\begin{pmatrix} 1 & 2 & 3 & 4 \\ 1 & 4 & 9 & 16 \\ 1 & 8 & 27 & 64 \\ 1 & 16 & 81 & 256 \\ 1 & 32 & 243 & 1024 \end{pmatrix}$ 
```

```
n=5
g<-matrix(0,nrow=n,ncol=4)
a<-rep(1:n,1)
a
```

```
## [1] 1 2 3 4 5
```

```
x<-c(1,2,3,4)
for (j in 1:n)
{g[j,]<-(x^j)}
g
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    2    3    4
## [2,]    1    4    9   16
## [3,]    1    8   27   64
## [4,]    1   16   81  256
## [5,]    1   32  243 1024
```

*#Absolutely ridiculous. I'm shocked that I actually did that.*

#### 4. Basic programming (10 points)

1. Let  $h(x, n) = 1 + x + x^2 + \dots + x^n = \sum_{i=0}^n x^i$ . Write an R program to calculate  $h(x, n)$  using a for loop. (5 points)

```
x<-2
h<-2
n<-5
x.fun<-rep(0,n)
for(i in 0:(n-1)) {
  x.fun[i]<-x^(i+1)
  y<-1+x+sum(x.fun)
}
x.fun
```

```
## [1] 4 8 16 32 0
```

```
y
```

```
## [1] 63
```

```
1+h+h**2+h**3+h**4+h**5
```

```
## [1] 63
```

```
sum(x^(0:n))
```

```
## [1] 63
```

```
#Here's an easy non-loop way to do it.
```

```
rowSums(outer(x, 0:n, "^"))
```

```
## [1] 63
```

```
#Here's another super cool way to do it.
```

1. If we list all the natural numbers below 10 that are multiples of 3 or 5, we get 3, 5, 6 and 9. The sum of these numbers is 23.

1. Find the sum of all the multiples of 3 or 5 below 1,000. (3, [euler1](#))

```
#Here's my best shot
```

```
x<-numeric(999)
```

```
#creating a vector that holds the values
```

```
for(i in 1:999) {
```

```
if (i %3 == 0 || i%5==0) {
```

```
x[i]<-i} else {x[i]<-0}}
```

```
x
```

```
## [1] 0 0 3 0 5 6 0 0 9 10 0 12 0 0 15 0 0
## [18] 18 0 20 21 0 0 24 25 0 27 0 0 30 0 0 33 0
## [35] 35 36 0 0 39 40 0 42 0 0 45 0 0 48 0 50 51
## [52] 0 0 54 55 0 57 0 0 60 0 0 63 0 65 66 0 0
## [69] 69 70 0 72 0 0 75 0 0 78 0 80 81 0 0 84 85
## [86] 0 87 0 0 90 0 0 93 0 95 96 0 0 99 100 0 102
## [103] 0 0 105 0 0 108 0 110 111 0 0 114 115 0 117 0 0
## [120] 120 0 0 123 0 125 126 0 0 129 130 0 132 0 0 135 0
## [137] 0 138 0 140 141 0 0 144 145 0 147 0 0 150 0 0 153
## [154] 0 155 156 0 0 159 160 0 162 0 0 165 0 0 168 0 170
## [171] 171 0 0 174 175 0 177 0 0 180 0 0 183 0 185 186 0
## [188] 0 189 190 0 192 0 0 195 0 0 198 0 200 201 0 0 204
## [205] 205 0 207 0 0 210 0 0 213 0 215 216 0 0 219 220 0
## [222] 222 0 0 225 0 0 228 0 230 231 0 0 234 235 0 237 0
## [239] 0 240 0 0 243 0 245 246 0 0 249 250 0 252 0 0 255
## [256] 0 0 258 0 260 261 0 0 264 265 0 267 0 0 270 0 0
## [273] 273 0 275 276 0 0 279 280 0 282 0 0 285 0 0 288 0
## [290] 290 291 0 0 294 295 0 297 0 0 300 0 0 303 0 305 306
## [307] 0 0 309 310 0 312 0 0 315 0 0 318 0 320 321 0 0
## [324] 324 325 0 327 0 0 330 0 0 333 0 335 336 0 0 339 340
## [341] 0 342 0 0 345 0 0 348 0 350 351 0 0 354 355 0 357
## [358] 0 0 360 0 0 363 0 365 366 0 0 369 370 0 372 0 0
## [375] 375 0 0 378 0 380 381 0 0 384 385 0 387 0 0 390 0
```

```
## [392] 0 393 0 395 396 0 0 399 400 0 402 0 0 405 0 0 408
## [409] 0 410 411 0 0 414 415 0 417 0 0 420 0 0 423 0 425
## [426] 426 0 0 429 430 0 432 0 0 435 0 0 438 0 440 441 0
## [443] 0 444 445 0 447 0 0 450 0 0 453 0 455 456 0 0 459
## [460] 460 0 462 0 0 465 0 0 468 0 470 471 0 0 474 475 0
## [477] 477 0 0 480 0 0 483 0 485 486 0 0 489 490 0 492 0
## [494] 0 495 0 0 498 0 500 501 0 0 504 505 0 507 0 0 510
## [511] 0 0 513 0 515 516 0 0 519 520 0 522 0 0 525 0 0
## [528] 528 0 530 531 0 0 534 535 0 537 0 0 540 0 0 543 0
## [545] 545 546 0 0 549 550 0 552 0 0 555 0 0 558 0 560 561
## [562] 0 0 564 565 0 567 0 0 570 0 0 573 0 575 576 0 0
## [579] 579 580 0 582 0 0 585 0 0 588 0 590 591 0 0 594 595
## [596] 0 597 0 0 600 0 0 603 0 605 606 0 0 609 610 0 612
## [613] 0 0 615 0 0 618 0 620 621 0 0 624 625 0 627 0 0
## [630] 630 0 0 633 0 635 636 0 0 639 640 0 642 0 0 645 0
## [647] 0 648 0 650 651 0 0 654 655 0 657 0 0 660 0 0 663
## [664] 0 665 666 0 0 669 670 0 672 0 0 675 0 0 678 0 680
## [681] 681 0 0 684 685 0 687 0 0 690 0 0 693 0 695 696 0
## [698] 0 699 700 0 702 0 0 705 0 0 708 0 710 711 0 0 714
## [715] 715 0 717 0 0 720 0 0 723 0 725 726 0 0 729 730 0
## [732] 732 0 0 735 0 0 738 0 740 741 0 0 744 745 0 747 0
## [749] 0 750 0 0 753 0 755 756 0 0 759 760 0 762 0 0 765
## [766] 0 0 768 0 770 771 0 0 774 775 0 777 0 0 780 0 0
## [783] 783 0 785 786 0 0 789 790 0 792 0 0 795 0 0 798 0
## [800] 800 801 0 0 804 805 0 807 0 0 810 0 0 813 0 815 816
## [817] 0 0 819 820 0 822 0 0 825 0 0 828 0 830 831 0 0
## [834] 834 835 0 837 0 0 840 0 0 843 0 845 846 0 0 849 850
## [851] 0 852 0 0 855 0 0 858 0 860 861 0 0 864 865 0 867
## [868] 0 0 870 0 0 873 0 875 876 0 0 879 880 0 882 0 0
## [885] 885 0 0 888 0 890 891 0 0 894 895 0 897 0 0 900 0
## [902] 0 903 0 905 906 0 0 909 910 0 912 0 0 915 0 0 918
## [919] 0 920 921 0 0 924 925 0 927 0 0 930 0 0 933 0 935
## [936] 936 0 0 939 940 0 942 0 0 945 0 0 948 0 950 951 0
## [953] 0 954 955 0 957 0 0 960 0 0 963 0 965 966 0 0 969
## [970] 970 0 972 0 0 975 0 0 978 0 980 981 0 0 984 985 0
## [987] 987 0 0 990 0 0 993 0 995 996 0 0 999
```

```
#this is a vector with all the multiples, with 0's in non-multiples
```

```
y<-x[x !=0]
```

```
#remove the zero's
```

```
y
```

```
## [1] 3 5 6 9 10 12 15 18 20 21 24 25 27 30 33 35 36
## [18] 39 40 42 45 48 50 51 54 55 57 60 63 65 66 69 70 72
## [35] 75 78 80 81 84 85 87 90 93 95 96 99 100 102 105 108 110
## [52] 111 114 115 117 120 123 125 126 129 130 132 135 138 140 141 144 145
## [69] 147 150 153 155 156 159 160 162 165 168 170 171 174 175 177 180 183
## [86] 185 186 189 190 192 195 198 200 201 204 205 207 210 213 215 216 219
## [103] 220 222 225 228 230 231 234 235 237 240 243 245 246 249 250 252 255
## [120] 258 260 261 264 265 267 270 273 275 276 279 280 282 285 288 290 291
## [137] 294 295 297 300 303 305 306 309 310 312 315 318 320 321 324 325 327
## [154] 330 333 335 336 339 340 342 345 348 350 351 354 355 357 360 363 365
## [171] 366 369 370 372 375 378 380 381 384 385 387 390 393 395 396 399 400
## [188] 402 405 408 410 411 414 415 417 420 423 425 426 429 430 432 435 438
```



```
## [205] 440 441 444 445 447 450 453 455 456 459 460 462 465 468 470 471 474
## [222] 475 477 480 483 485 486 489 490 492 495 498 500 501 504 505 507 510
## [239] 513 515 516 519 520 522 525 528 530 531 534 535 537 540 543 545 546
## [256] 549 550 552 555 558 560 561 564 565 567 570 573 575 576 579 580 582
## [273] 585 588 590 591 594 595 597 600 603 605 606 609 610 612 615 618 620
## [290] 621 624 625 627 630 633 635 636 639 640 642 645 648 650 651 654 655
## [307] 657 660 663 665 666 669 670 672 675 678 680 681 684 685 687 690 693
## [324] 695 696 699 700 702 705 708 710 711 714 715 717 720 723 725 726 729
## [341] 730 732 735 738 740 741 744 745 747 750 753 755 756 759 760 762 765
## [358] 768 770 771 774 775 777 780 783 785 786 789 790 792 795 798 800 801
## [375] 804 805 807 810 813 815 816 819 820 822 825 828 830 831 834 835 837
## [392] 840 843 845 846 849 850 852 855 858 860 861 864 865 867 870 873 875
## [409] 876 879 880 882 885 888 890 891 894 895 897 900 903 905 906 909 910
## [426] 912 915 918 920 921 924 925 927 930 933 935 936 939 940 942 945 948
## [443] 950 951 954 955 957 960 963 965 966 969 970 972 975 978 980 981 984
## [460] 985 987 990 993 995 996 999
```

```
#this is a vector with just the multiples
sum(y)
```

```
## [1] 233168
```

```
#The answer is 233168
```

1. Find the sum of all the multiples of 4 or 7 below 1,000,000. (2)

```
#Same deal as before, just different multiples.
a<-numeric(999999)
for(i in 1:999999){
  if (i %%4 ==0 || i%%7==0){
    a[i]<-i} else {a[i]<-0}
  }
head(a)
```

```
## [1] 0 0 0 4 0 0
```

```
b<-a[a !=0]
head(b)
```

```
## [1] 4 7 8 12 14 16
```

```
sum(b)
```

```
## [1] 178571071431
```

```
#answer is 178571071431
#You probably won't like the manual entry and deletion of 0 terms...
#... but this is the way I found that works.
```

1. Each new term in the Fibonacci sequence is generated by adding the previous two terms. By starting with 1 and 2, the first 10 terms will be (1, 2, 3, 5, 8, 13, 21, 34, 55, 89). Write an R program to calculate the sum of the first 15 even-valued terms. (5 bonus points, [euler2](#))

*#Going to try to determine the 15th even value term first*

```
bin <- 45
fibvals <- numeric(bin)
fibvals[1] <- 1
fibvals[2] <- 1
for (i in 3:bin) {
  fibvals[i] <- fibvals[i-1]+fibvals[i-2]
}
fibvals
```

```
## [1]      1      1      2      3      5      8
## [7]     13     21     34     55     89    144
## [13]    233    377    610    987   1597   2584
## [19]   4181   6765  10946  17711  28657  46368
## [25]  75025 121393 196418 317811 514229 832040
## [31] 1346269 2178309 3524578 5702887 9227465 14930352
## [37] 24157817 39088169 63245986 102334155 165580141 267914296
## [43] 433494437 701408733 1134903170
```

*#Here are the first 45 values of the fibonacci sequence*

*#(I specified bin with some prior knowledge that fib. sequence produces...*

*#...an even number every 3rd term. Might be cheating a little)*

```
c<-numeric(15)
for(i in seq(fibvals)) {
  if(fibvals[i]%2==0) {
    c[i]<-fibvals[i]} else{c[i]<-0}
  }
}
head(c)
```

```
## [1] 0 0 2 0 0 8
```

```
c
```

```
## [1]      0      0      2      0      0      8
## [7]      0      0     34      0      0    144
## [13]      0      0    610      0      0   2584
## [19]      0      0   10946      0      0  46368
## [25]      0      0   196418      0      0 832040
## [31]      0      0   3524578      0      0 14930352
## [37]      0      0  63245986      0      0 267914296
## [43]      0      0 1134903170
```

```
ci<-c[c !=0]
length(ci)
```

```
## [1] 15
```

```
ci
```

```
## [1]      2      8     34     144     610    2584
## [7]   10946   46368  196418  832040 3524578 14930352
## [13] 63245986 267914296 1134903170
```

```
sum(ci)
```

```
## [1] 1485607536
```

```
#The answer is 1485607536.
```

```
#Definitely not the most efficient way of doing it, but given that ...
```

```
#... I have almost no prior experience with loops, I'm happy with it.
```

Some problems taken or inspired by projecteuler.